

Polistes dominula genome project

Daniel Standage

Volker Brendel

Amy Toth

November 13, 2014

1 Overview

This documentation is a record of our work for the *Polistes dominula* genome project. It was created to 1) serve as full disclosure of all of the methods, commands, and software used to produce the reported results, and 2) facilitate anonymous replication of those results.

1.1 Data access

Raw instrument data and final data outputs are stored in the [iPlant Data Store](#) under the path `/iplant/home/standage/Polistes_dominula/`. All file and directory paths provided in this documentation are relative to that root path, which for the remainder of the documentation will be designated the **Pdom Data Store**.

1.2 Using this documentation

This project is divided into several sections, with each section focusing on a single analysis or small group of related analyses. Each section has a dedicated directory containing code and documentation specific to that section. These resources can be browsed or downloaded at [GitHub](#).

- A `README.md` file (in Markdown format) is included for each section, which provides a prose description of what each set of commands is doing. This file is intended to facilitate interactive replication of results: typing or pasting the commands into the terminal and executing them manually to produce the output. (Note: a single PDF document containing all documentation was produced by concatenating all of the various README files into a single Markdown file and converting to PDF format.)
- Each section also contains a `Makefile` file which includes the same commands as the README file, though without the commentary and in slightly different syntax. The purpose of these files is to facilitate automated replication of each analysis in batch mode. To execute this procedure for a particular analysis, simply change to that directory and execute `make` on the command line.
- Most sections also include additional supplementary files, such as source code, graphics, or configuration files necessary for replicating the results. The purpose of each supplemental file should be clear from the documentation.

If you encounter any problems using this documentation or its associated files, please open a ticket with the [Pdom Genome Project issue tracker](#).

1.3 Authors

- [Daniel Standage](#); Indiana University
- [Volker Brendel](#); Indiana University
- [Amy Toth](#), principal investigator; Iowa State University

2 Genome size estimation

[Jellyfish](#) version 2.1.3 was used to count k -mer distributions in the raw genomic short read data. The k -mer coverage C_k was determined for several values of k : 17, 21, 25, and 29. A linear model of C_k as a function of k was fit to compute the estimated nucleotide coverage $C = C_1$ and genome size. The k -mer histogram files have been deposited in the Pdom Data Store at [r1.2/genome-size-est/](#).

2.1 Procedure (interactive)

First, designate the number of available processors. This will run multiple jobs/threads at once to speed up computations. For a laptop or a desktop, this will usually be 4, 8, or 16. For server or HPC hardware, you may have as many as 32 to 64 processors at your disposal.

```
NumThreads=16
```

Next, download short reads using [iRODS](#) and decompress.

```
iget -Vr /iplant/home/standage/Polistes_dominula/sequence/genome
ls genome/*.gz | parallel --gnu --jobs $NumThreads gunzip
```

Then, count k -mers and produce k -mer frequency histograms.

```
FastqFiles=$(ls genome/*.fq)
for k in 17 21 25 29
do
    jellyfish count -m $k -s 100M -t $NumThreads -C -o pdom-${k}mers.jf $FastqFiles
    jellyfish histo pdom-${k}mers.jf > pdom-${k}mers.hist
done
```

Finally, estimate k -mer coverage, genome coverage, and genome size.

```
./size-coverage-estimate.R
```

Clean up huge data files.

```
rm -r genome/*.fq *.jf
```

2.2 Procedure (automated)

The same procedure can also be run in batch mode using the following commands (in the `genome-size` directory).

```
make NumThreads=16
make clean
```

2.3 References

- Marçais G, Kingsford C (2011) A fast, lock-free approach for efficient parallel counting of occurrences of k -mers. *Bioinformatics* 27:764-70, [doi:10.1093/bioinformatics](#).

3 Genome assembly

Raw DNA-Seq reads were groomed using [Trimmomatic](#) version [0.22](#), and the groomed reads were then assembled using [AllPaths-LG](#) version [43216](#). The final assembly file has been deposited in the Pdom Data Store at [r1.2/genome-assembly/pdom-scaffolds-unmasked-r1.2.fa.gz](#).

3.1 Procedure (interactive)

3.1.1 Short read quality control

First, designate the number of available processors to speed up Trimmomatic's computations. Also, provide the path of the `trimmomatic-0.22.jar` file contained in the Trimmomatic source code distribution.

```
NumThreads=16
TrimJar=/usr/local/src/Trimmomatic-0.22/trimmomatic-0.22.jar
PdomData=/iplant/home/standage/Polistes_dominula
```

Now for the processing. We apply the following filters to each read pair.

- remove adapter contamination
- remove any nucleotides at either end of the read whose quality score is below 3
- trim the read once the average quality in a 5bp sliding window falls below 20
- discard any reads which, after processing, fall below the length threshold (40bp for 100bp reads, 26bp for 35bp reads)

```
for sample in 200bp 500bp 1kb 3kb 8kb
do
  iget ${PdomData}/sequence/genome/pdom-gdnaseq-${sample}-1.fq.gz
  iget ${PdomData}/sequence/genome/pdom-gdnaseq-${sample}-2.fq.gz
  ./run-trim.sh $sample $TrimJar $NumThreads
done
```

3.1.2 Assembly with AllPaths-LG

First, prepare a working directory for the assembly.

```
mkdir -p assembly/Polistes_dominula/data-trim
mv pdom-gdnaseq-*-trim-?.fq assembly/Polistes_dominula/data-trim/.
```

Next, convert the input files into the internal format required by AllPaths-LG.

```
PrepareAllPathsInputs.pl \
  DATA_DIR=assembly/Polistes_dominula/data-trim \
  PLOIDY=2
```

Then, execute the assembly procedure.

```
RunAllPathsLG PRE=assembly \
  REFERENCE_NAME=Polistes_dominula \
  DATA_SUBDIR=data-trim \
  RUN=run01 \
  TARGETS=standard
```

Finally, assign official project IDs to the scaffolds, compress, and clean up intermediate data files.

```
./scaff-ids.pl PdomSCFr1.2- \  
  < $PRE/Polistes_dominula/data-trim/run01/ASSEMBLIES/test/final.assembly.fasta \  
  > pdom-scaffolds-unmasked-r1.2.fa  
gzip pdom-scaffolds-unmasked-r1.2.fa  
rm -rf assembly pdom-gdnaseq*.fq*
```

3.2 Procedure (automated)

The same procedure can also be run in batch mode using the following commands (in the `genome-assembly` directory).

```
make NumThreads=16 \  
  TrimJar=/usr/local/src/Trimmomatic-0.22/trimmomatic-0.22.jar  
make clean
```

3.3 References

- Lohse M, Bolger AM, Nagel A, Fernie AR, Lunn JE, Stitt M, Usadel B (2012) RobiNA: a user-friendly, integrated software solution for RNA-Seq-based transcriptomics. *Nucleic Acids Research*, 40:W622-7, doi:10.1093/nar/gks540.
- Gnerre S, MacCallum I, Przybylski D, Ribeiro F, Burton J, Walker B, Sharpe T, Hall G, Shea T, Sykes S, Berlin A, Aird D, Costello M, Daza R, Williams L, Nicol R, Gnirke A, Nusbaum C, Lander ES, Jaffe DB (2010) High-quality draft assemblies of mammalian genomes from massively parallel sequence data. *Proceedings of the National Academy of Sciences USA*, 108(4):1513-1518, doi:10.1073/pnas.1017351108.

4 Transcript assembly

Raw DNA-Seq reads were groomed using [Trimmomatic](#) version 0.22, and the groomed reads were then assembled using [Trinity](#) version r20131110. Then the assembled transcripts were processed with [mRNAmarkup](#) version 10-3-2013 to remove contaminants and correct erroneously assembled chimeric transcripts. The cleaned and annotated transcripts have been deposited in the Pdom Data Store at `r1.2/transcript-assembly/`.

4.1 Procedure (interactive)

4.1.1 Short read quality control

First, designate the number of available processors to speed up Trimmomatic's computations. Also, provide the path of the `trimmomatic-0.22.jar` file contained in the Trimmomatic source code distribution.

```
NumThreads=16  
TrimJar=/usr/local/src/Trimmomatic-0.22/trimmomatic-0.22.jar  
PdomData=/iplant/home/standage/Polistes_dominula
```

Now for the processing. We apply the following filters to each read pair.

- remove adapter contamination

- remove any nucleotides at either end of the read whose quality score is below 3
- trim the read once the average quality in a 5bp sliding window falls below 20
- discard any reads which, after processing, fall below the 40bp length threshold

```
for caste in q w
do
  for rep in {1..6}
  do
    sample=${caste}${rep}
    iget ${PdomData}/sequence/transcriptome/pdom-rnaseq-${sample}-1.fq.gz
    iget ${PdomData}/sequence/transcriptome/pdom-rnaseq-${sample}-2.fq.gz
    ./run-trim.sh $sample $TrimJar $NumThreads
  done
done
```

4.1.2 Assembly with Trinity

Trinity requires a single input file—or a pair of input files for paired-end data. We need to combine all of the data into a single pair of files.

```
cat pdom-rnaseq-*-trim-1.fq > pdom-rnaseq-all-trim-1.fq
cat pdom-rnaseq-*-trim-2.fq > pdom-rnaseq-all-trim-2.fq
```

We'll then execute the Trinity assembler using the `--CuffFly` reconstruction algorithm.

```
Trinity.pl --seqType fq \
  --JM 100G \
  --bflyHeapSpaceMax 50G \
  --output pdom-trinity \
  --CPU $NumThreads \
  --left pdom-rnaseq-all-trim-1.fq \
  --right pdom-rnaseq-all-trim-2.fq \
  --full_cleanup \
  --jaccard_clip \
  --CuffFly
```

4.1.3 Post-processing with mRNAmakup

Contaminant, reference protein, and miRNA databases were collected as described in the mRNAmakup documentation (db/OREADME and db/OREADME-hy). The mRNAmakup procedure was then run on the Trinity output. Be sure to edit the mRNAmakup.conf file with the correct paths to the databases.

```
mRNAmakup -c mRNAmakup.conf \
  -i pdom-trinity/Trinity.fasta \
  -o output-mRNAmakup
```

4.1.4 Potential *Polistes*-specific genes

Polistes metricus and *P. canadensis* transcript assemblies were groomed and annotated using the same mRNAmakup procedure. All three transcript sets had a substantial number of TSAs that could not be

annotated by mRNAMarkup. The longest open reading frame translations of at least 80 amino acids derived from these unmatched TSAs were pairwise compared with BLASTp to detect any *Polistes*-conserved and species-specific genes. This procedure relies on a workflow and several scripts available in the supplemental/directory of the documentation repository.

The workflow to find the *Polistes*-conserved transcripts with no external protein matches can be executed with this command. The workflow depends on the `dnatopro` program (part of the VBTools utilities included in the mRNAMarkup distribution) and the NCBI BLAST+ suite.

```
./venn.make BLASTTHREADS=$NumThreads all
```

Transcripts associated with potential *Polistes*-specific genes will be placed in the `pdom-tsa-r1.2-unmatched-pep.fa` file.

4.2 Procedure (automated)

The same procedure can also be run in batch mode using the following commands (in the `transcript-assembly` directory). Note that this procedure does not automate the mRNAMarkup analysis. After producing the Trinity assembly, it retrieves previously computed mRNAMarkup results for the clade-specific gene analysis.

```
make NumThreads=16 \
  TrimJar=/usr/local/src/Trimmomatic-0.22/trimmomatic-0.22.jar
make clean
```

4.3 References

- Lohse M, Bolger AM, Nagel A, Fernie AR, Lunn JE, Stitt M, Usadel B (2012) RobiNA: a user-friendly, integrated software solution for RNA-Seq-based transcriptomics. *Nucleic Acids Research*, 40:W622-7, doi:10.1093/nar/gks540.
- Grabherr MG, Haas BJ, Yassour M, Levin JZ, Thompson DA, Amit I, Adiconis X, Fan L, Raychowdhury R, Zeng Q, Chen Z, Mauceli E, Hacohen N, Gnirke A, Rhind N, di Palma F, Birren BW, Nusbaum C, Lindblad-Toh K, Friedman N, Regev A (2011) Full-length transcriptome assembly from RNA-seq data without a reference genome. *Nature Biotechnology*, 29(7):644-52, doi:10.1038/nbt.1883.

5 Repeat masking

The genome assembly was screened for known repetitive elements using RepeatMasker version [open-4.0.5](#) and [Repbase](#) version 20140131. After masking repeats identified by RepeatMasker, the assembly was screened for additional repeats using [Tallymer](#) version 1.5.2. To discriminate *bona fide* repetitive elements from genes occurring in high copy number in the genome, all repeats identified by Tallymer were subjected to a BLASTX search against a database of Hexapod proteins. Any repeats with matches in the database and e-values < 1e-5 were discarded as probable high copy number genes, while the rest were used to mask the genome. The final masked sequence has been deposited in the Pdom data store at `r1.2/genome-assembly/pdom-scaffolds-masked-r1.2.fa.gz`.

5.1 Procedure (interactive)

First, download the unmasked genome sequence.

```
PdomData=/iplant/home/standage/Polistes_dominula
iget ${PdomData}/r1.2/genome-assembly/pdom-scaffolds-unmasked-r1.2.fa.gz
gunzip pdom-scaffolds-unmasked-r1.2.fa.gz
```

5.1.1 Screening with RepeatMasker

Next, identify known repeats with RepeatMasker. By default, RepeatMasker produces soft-masked (lower-case) sequences, so we need to post-process the output to hard mask (N) the sequence.

```
NumThreads=16
GCContent=30.77
RepeatMasker -species insects -parallel $NumThreads -gc $GCContent \
  -frag 4000000 -lcambig -xsmall -gff \
  pdom-scaffolds-unmasked-r1.2.fa \
  > rm.log 2>&1
python lc2n.py < pdom-scaffolds-unmasked-r1.2.fa.masked \
  > pdom-rm-masked.fa
```

5.1.2 Screening with Tallymer

Then, do additional *k*-mer based screening for repetitive elements using Tallymer (procedure published by [Dan Bolser](#)).

```
gt suffixerator -v \
  -db $IDX \
  -indexname $IDX \
  -tis -suf -lcp -des -ssp -sds -dna \
  > suffixerator.log 2>&1

gt tallymer occratio -v \
  -minmersize 10 \
  -maxmersize 45 \
  -output unique nonunique nonuniquemulti total relative \
  -esa $IDX \
  > pdom.occratio.10.45.dump

gt tallymer mkindex -v \
  -mersize 19 \
  -minocc 50 \
  -esa $IDX \
  -counts -pl \
  -indexname pdom.idx.19.50 \
  > mkindex.log 2>&1

gt tallymer search -v \
  -output qseqnum qpos counts \
  -tyr pdom.idx.19.50 \
  -q $IDX \
  > pdom.repeats.19.50.tmer \
  2> tallymer.search.log
```

```
tallymer2gff3.plx -k 19 -s $IDX \
                  pdom.repeats.19.50.tmer \
                  > pdom.repeats.19.50.gff3

gff2fasta.plx -s pdom-rm-masked.fa \
              -f pdom.repeats.19.50.gff3 \
              > pdom.repeats.19.50.fa
```

Do a BLASTx search of repeats found by Tallymer vs known hexapod proteins, and parse out those with hits using [MuSeqBox](#).

```
curl 'http://www.uniprot.org/uniprot/?query=taxonomy%3a6960&force=yes&format=fasta' \
    > hexapoda.fa
makeblastdb -in hexapoda.fa -dbtype prot -parse_seqids
blastx -query pdom.repeats.19.50.fa -db hexapoda.fa \
       -num_alignments 10 -evalue 1e-5 -num_threads 64 \
       -out pdom.repeats.19.50.blastx \
       > pdom.repeats.19.50.log 2>&1

MuSeqBox -i pdom.repeats.19.50.blastx -L 100 \
  | cut -f 1 -d ' ' | sort | uniq \
  | perl -ne 'm/(PdomSCFr1.2-\d+)-\d+\/(\d+)-(\d+)/ and print "$1\t$2\t$3\n"' \
  > pdom.repeats.19.50.hexapodhits.txt
```

Finally, hard mask the Tallymer repeats, excluding any that match Hexapod proteins as probably high-copy-number genes.

```
mask.pl pdom.repeats.19.50.gff3 \
        pdom.repeats.19.50.hexapodhits.txt \
        pdom-rm-masked.fa \
        > pdom-scaffolds-masked-r1.2.fa
gzip pdom-scaffolds-masked-r1.2.fa
```

5.2 Procedure (automated)

The same procedure can also be run in batch mode using the following commands (in the `repeat-masking` directory).

```
make NumThreads=16 GCCContent=30.77
make clean
```

5.3 References

- Smit AFA, Hubley R, Green P. RepeatMasker Open-3.0. 1996-2010 <http://www.repeatmasker.org>.
- Jurka J, Kapitonov VV, Pavlicek A, Klonowski P, Kohany O, Walichiewicz J (2005) Repbase Update, a database of eukaryotic repetitive elements. *Cytogenetic and Genome Research*, **110**(1-4):462-467, doi:10.1159/000084979.
- Kurtz S, Narechania A, Stein JC, Ware D (2009) A new method to compute *K*-mer frequencies and its application to annotate large repetitive plant genomes. *BMC Genomics*, **9**:517, doi:10.1186/1471-2164-9-517.

6 Transcript alignment

Assembled and groomed *P. dominula* transcripts were spliced aligned at high stringency to the repeat-masked genome sequence using [GeneSeqer](#) version 2-26-2014. *Polistes canadensis* and *P. metricus* transcripts were also aligned with GeneSeqer using less stringent parameters. The alignments have been deposited in the Pdom Data Store at [r1.2/transcript-alignment](#).

6.1 Procedure (interactive)

First, set the `GSQDir` variable to the path of the GeneSeqer source code.

```
GSQDir=/usr/local/src/GENESEQER
```

Next, align *P. dominula* TSAs with stringent parameters.

```
# Download the assembled and masked genome sequence
PdomData=/iplant/home/standage/Polistes_dominula
iget ${PdomData}/r1.2/genome-assembly/pdom-scaffolds-masked-r1.2.fa.gz
gunzip pdom-scaffolds-masked-r1.2.fa.gz

# P. dominula alignments
iget ${PdomData}/r1.2/transcript-assembly/pdom-tsa-r1.2.fa.gz
gunzip pdom-tsa-r1.2.fa.gz
MakeArray pdom-tsa-r1.2.fa
GeneSeqerL -s Arabidopsis \
           -L pdom-scaffolds-masked-r1.2.fa \
           -D pdom-tsa-r1.2.fa.gz \
           -O pdom-tsa-r1.2-masked.gsq \
           -p $GSQDir/data/prmfileHQ \
           -x 30 -y 45 -z 60 -w 0.8 -m 1000000 \
           > pdom-tsa-r1.2-masked-gsq.log 2>&1
```

Now the *P. canadensis* and *P. metricus* TSAs with less stringent parameters.

```
# P. metricus alignments
iget ${PdomData}/r1.2/transcript-assembly/pmet-tsa-r1.2.fa.gz
gunzip pmet-tsa-r1.2.fa.gz
MakeArray pmet-tsa-r1.2.fa
GeneSeqerL -s Arabidopsis \
           -L pdom-scaffolds-masked-r1.2.fa \
           -d pmet-tsa-r1.2.fa \
           -O pmet-tsa-r1.2-masked.gsq \
           -p $GSQDir/data/prmfile \
           -x 16 -y 24 -z 48 -w 0.8 -m 1000000 \
           > pmet-tsa-r1.2-masked-gsq.log 2>&1

# P. canadensis alignments
iget ${PdomData}/r1.2/transcript-assembly/pcan-tsa.fa.gz
gunzip pcan-tsa.fa.gz
MakeArray pcan-tsa.fa
GeneSeqerL -s Arabidopsis \
           -L pdom-scaffolds-masked-r1.2.fa \
```

```
-d pcan-tsa.fa \
-O pcan-tsa-masked.gsq \
-p $GSQDIR/data/prmfile \
-x 16 -y 24 -z 48 -w 0.8 -m 1000000 \
> pcan-tsa-masked-gsq.log 2>&1
```

Lastly, convert all of the alignments into GFF3 format, filtering out alignments with GeneSeqer similarity or coverage scores < 0.8.

```
./gsq2makergff3.py < pdom-tsa-r1.2-masked.gsq > pdom-tsa-r1.2-masked.gff3
./gsq2makergff3.py < pmet-tsa-r1.2-masked.gsq > pmet-tsa-r1.2-masked.gff3
./gsq2makergff3.py < pcan-tsa-r1.2-masked.gsq > pcan-tsa-masked.gff3
```

6.2 Procedure (automated)

The same procedure can also be run in batch mode using the following commands (in the `transcript-alignment` directory).

```
make GSQDir=/usr/local/src/GENESEQER
```

6.3 References

- **Brendel V, Xing L, Zhu W** (2004) Gene structure prediction from consensus spliced alignment of multiple ESTs matching the same genomic locus. *Bioinformatics*, **20**:1157-1169, [doi:10.1093/bioinformatics/bth058](https://doi.org/10.1093/bioinformatics/bth058).
- **Berens *et al***, in preparation.
- **Ferreira P, Patalano S, Chauhan R, Ffrench-Constant R, Gabaldon T, Guigo R, Sumner S** (2013) Transcriptome analyses of primitively eusocial wasps reveal novel insights into the evolution of sociality and the origin of alternative phenotypes. *Genome Biology*, **14**(2):R20, [doi:10.1186/gb-2013-14-2-r20](https://doi.org/10.1186/gb-2013-14-2-r20).

7 Reference protein alignment

Reference protein sequences from *Apis mellifera* ([OGS 3.2](#) & [NCBI GNOMON](#)) and *Drosophila melanogaster* ([FlyBase r5.55](#)) were splice-aligned to the genome using [GenomeThreader](#) version 1.6.0. The alignments have been deposited in the Pdom Data Store at `r1.2/protein-alignment`.

7.1 Procedure (interactive)

First, download the assembled and masked genome sequence.

```
PdomData=/iplant/home/standage/Polistes_dominula
iget ${PdomData}/r1.2/genome-assembly/pdom-scaffolds-masked-r1.2.fa.gz
gunzip pdom-scaffolds-masked-r1.2.fa.gz
```

Next, download the reference proteins.

```

BeeData=http://hymenopteragenome.org/beebase/sites/hymenopteragenome.org.beebase
curl ${BeeData}/files/data/consortium_data/amel_OGSv3.2_pep.fa.gz \
| zcat \
| sed 's/gnl|Amel_4.5|//g' \
> amel-ogs-prot.fa
curl ftp://ftp.ncbi.nih.gov/genomes/Apis_mellifera/protein/protein.fa.gz \
| zcat \
| perl -ne 's/>gi\|(\d+)\|ref\|([^\|]+)\|S*/>$1 $2/; print' \
> amel-ncbi-prot.fa
FlyData=ftp://ftp.flybase.net/releases/FB2014_01/dmel_r5.55
curl ${FlyData}/fasta/dmel-all-translation-r5.55.fasta.gz \
| zcat \
> dmel-flybase-prot.fa

```

Then perform the alignments.

```

# The $GTHBIN variable should point to the bin directory in the GenomeThreader
# distribution.
export BSSMDIR=$GTHBIN/bssm
export GTHDATADIR=$GTHBIN/gthdata

for prot in amel-ogs amel-ncbi dmel-flybase
do
    $GTHBIN/gth -genomic pdom-scaffolds-masked-r1.2.fa \
    -protein ${prot}-prot.fa \
    -species arabidopsis \
    -gcmaxgapwidth 20000 \
    -gcmincoverage 25 \
    -prhdist 6 \
    -prminmatchlen 18 \
    -prseedlength 6 \
    -o ${prot}-prot-masked.gth \
    -force \
    > ${prot}-prot-masked.log 2>&1
done

```

Finally, convert the GenomeThreader alignments into GFF3 format, filtering out alignments with similarity or coverage scores < 0.5.

```

./gth2makergff3.py < amel-ogs-prot-masked.gth > amel-ogs-prot-masked.gff3
./gth2makergff3.py < amel-ncbi-prot-masked.gth > amel-ncbi-prot-masked.gff3
./gth2makergff3.py < dmel-flybase-prot-masked.gth > dmel-flybase-prot-masked.gff3

```

7.2 Procedure (automated)

The same procedure can also be run in batch mode using the following commands (in the `protein-alignment` directory).

```
make GTHbin=/usr/local/bin/GTH
```

7.3 References

- Gremme G, Brendel V, Sparks ME, Kurtz S (2005) Engineering a software tool for gene structure prediction in higher organisms. *Information and Software Technology*, **47**(15):965-978, doi:10.1016/j.infsof.2005.09.005.

8 Genome annotation

The [Maker annotation pipeline](#) version 2.31.6 was used to annotate protein-coding genes and tRNA genes in the *P. dominula* genome. The final, cleaned up annotation file and corresponding sequence files have been deposited in the Pdom Data Store at [r1.2/genome-annotation/](#).

8.1 Gene predictor training

Approximately 3000 genes were selected from annotations whose reliability was confirmed by strict measures of conservation with several other Hymenopteran species. These genes were then used to train 3 ab initio gene predictors (SNAP, Augustus, and GeneMark) for use with the Maker pipeline. For Augustus and SNAP, the procedures documented in their respective source code distributions were followed. For GeneMark, a model built by the self-training version of GeneMark was used with Maker. All models have been deposited in the Pdom Data Store at [r1.2/genome-annotation/models](#).

8.2 Procedure (interactive)

8.2.1 Data files

First, there are *a lot* of data files to download. Make sure to set the `AUGUSTUS_CONFIG_DIR` variable appropriately.

```
PdomData=/iplant/home/standage/Polistes_dominula
AUGUSTUS_CONFIG_DIR=/usr/local/src/augustus/config
export AUGUSTUS_CONFIG_DIR

# Genome sequence
iget ${PdomData}/genome-assembly/pdom-scaffolds-masked-r1.2.fa.gz
gunzip pdom-scaffolds-masked-r1.2.fa.gz

# Transcript alignments
iget ${PdomData}/r1.2/transcript-alignment/pdom-r1.2-trans-align.tar.gz
tar -xzf pdom-r1.2-trans-align.tar.gz
mv pdom-r1.2-trans-align/*.gff3 .

# Protein alignments
iget ${PdomData}/r1.2/protein-alignment/pdom-r1.2-refrprot-align.tar.gz
tar -xzf pdom-r1.2-refrprot-align.tar.gz
mv pdom-r1.2-refrprot-align/*.gff3 .

# Highly conserved and manually curated annotations
iget ${PdomData}/r1.2/genome-annotation/pdom-viga-improved.gff3.gz
iget ${PdomData}/r1.2/genome-annotation/pdom-yrgate.gff3.gz
gunzip pdom-viga-improved.gff3.gz
```

```
gunzip pdom-yrgate.gff3.gz

# Download species-specific parameter settings for ab initio gene predictors
iget ${PdomData}/r1.2/genome-annotation/models/pdom.snap.hmm
iget ${PdomData}/r1.2/genome-annotation/models/pdom.genemark.mod
iget ${PdomData}/r1.2/genome-annotation/models/pdom.augustus.tar.gz
tar -xzf pdom.augustus.tar.gz
mv pdom ${AUGUSTUS_CONFIG_DIR}/species
```

8.2.2 Configuration files

Next, we need to create the Maker control files. If all of the supplementary programs are in the path, the `maker_exe.ctl` file will be populated automatically. If not, you will need to manually fill it in with the location of all the programs.

```
maker -CTL
# Replace the empty `maker_opts.ctl` file with our file,
# which has all the values filled in.
cp pdom_maker_opts.ctl maker_opts.ctl
```

8.2.3 Annotation with Maker

After all of the data files and control files are in place, running Maker is trivial.

```
NumThreads=16
maker -genome pdom-scaffolds-masked-r1.2.fa \
      -fix_nucleotides \
      -nodatastore \
      -RM_off \
      -cpus $NumThreads \
      -base pdom \
      > pdom.maker.log 2>&1
```

8.2.4 Post-processing

The feature IDs created internally by Maker are pretty unwieldy, so we use a few scripts to clean them up.

```
# Merge and clean up data
gff3_merge -o pdom-annot-p1.2-raw.gff3 pdom.maker.output/pdom_master_datastore_index.log
bash clean.sh pdom-annot-p1.2-raw.gff3 > pdom-annot-p1.2-cleaned.gff3
# Fixes strange output artifact for some VIGA-based predictions
perl -n fix.pl < pdom-annot-p1.2-cleaned.gff3 > pdom-annot-p1.2-fixed.gff3

# Create a polished GFF3 file with official IDs for the annotations and proper
# ##sequence-region pragmas.
gt gff3 -retainids -sort -tidy pdom-annot-p1.2-fixed.gff3 2> gt.log \
  | python annot-ids.py --idfmt='Pdom%sr1.2-%05lu' -n --rnamap=rnaids.txt --dbxref=MAKER - \
  | python fix-regions.py pdom-scaffolds-masked-r1.2.fa > pdom-annot-r1.2.gff3

# Create transcript and protein sequence files with proper feature IDs
cat pdom.maker.output/pdom_datastore/PdomSCFr1.2-*/PdomSCFr1%2E2-*.maker.transcripts.fasta \
```

```

| python seq-ids.py rnaids.txt > pdom-annot-r1.2-transcripts.fasta
cat pdom.maker.output/pdom_datastore/PdomSCFr1.2-*/PdomSCFr1%2E2-*.maker.proteins.fasta \
| python seq-ids.py rnaids.txt > pdom-annot-r1.2-proteins.fasta

```

8.2.5 Functional annotation by BLASTp search

Translation products of all gene models predicted by Maker were searched against all animal proteins in the NCBI nr database. This BLAST search provides a preliminary functional annotation for the gene models, but also identified gene models without any matches to known proteins.

```

# NCBI nr database installed using update_blastdb.pl script
# distributed with BLAST.
blastdb_aliastool -gilist anm.p.gil -dbtype prot -db nr -out anm

# BLAST search
blastp -db anm \
    -query pdom-annot-r1.2-proteins.fa \
    -evalue 1e-4 \
    -num_threads $NumThreads \
    -outfmt 5 \
    -out pdom-r1.2-vs-anm-blastp.xml

```

Using the NCBI Taxonomy data, we can examine the best BLAST hit for each gene model and tally these up according to the family of the best hit.

```

# Download NCBI taxonomy data
mkdir taxonomy
cd taxonomy
curl -O ftp://ftp.ncbi.nih.gov/pub/taxonomy/taxdump.tar.gz
tar xzf taxdump.tar.gz
cd -

perl blastxml-to-besthit.pl taxonomy/names.dmp \
    < pdom-r1.2-vs-anm-blastp.xml \
    > pdom-annot-r1.2-besthit-gis.txt

# Breakdown of hits to Hymenoptera proteins, by family
./taxtrav --rank family \
    --filter 'order=Hymenoptera' \
    --nodes taxonomy/nodes.dmp \
    --names taxonomy/names.dmp \
    <(cut -f 2 pdom-annot-r1.2-besthit-gis.txt)
    > pdom-hym-hits-tax.csv
cut -f 3 -d , pdom-hym-hits-tax.csv | sort | uniq -c | sort -rn

# Number of hits to non-Hymenoptera proteins
./taxtrav --rank family \
    --filter 'order=Hymenoptera' \
    --nodes taxonomy/nodes.dmp \
    --names taxonomy/names.dmp \
    --complement \
    <(cut -f 2 pdom-annot-r1.2-besthit-gis.txt)

```

```
> pdom-nothym-hits-tax.csv
wc -l pdom-nothym-hits-tax.csv
```

Finally, let's pull out the gene models with no matches in the database. Unmatched gene models tend to be short, but we can also pull out gene models whose length (in amino acids) exceeds the median length for all gene models and set these aside as for further examination.

```
comm -13 <(cut -f 1 pdom-annot-r1.2-besthit-gis.txt | sort) \
    <(perl -ne 'm/>(\S+)/ and print "$1\n"' < pdom-annot-r1.2-proteins.fa) \
    > pdom-r1.2-no-anm-hits.ids
./select-seq pdom-r1.2-no-anm-hits.ids \
    pdom-annot-r1.2-proteins.fa \
    > pdom-r1.2-no-anm-hits.fa
bash fasta-median-plus.sh pdom-annot-r1.2-proteins.fa \
    pdom-r1.2-no-anm-hits.fa \
    > pdom-annot-r1.2-long-unmatched.fa
```

8.3 Procedure (automated)

The same procedure can also be run in batch mode using the following commands (in the `genome-annotation` directory). Keep in mind that unless all of the supplemental programs are in your path, you will still need to run `maker -CTL` and edit the `maker_exe.ctl` file manually.

```
make NumThreads=16
```

8.4 References

- Cantarel BL, Korf I, Robb SMC, Parra G, Ross E, Moore B, Holt C, Alvarado AS, Yandell M (2008) MAKER: An easy-to-use annotation pipeline designed for emerging model organism genomes. *Genome Research*, 18:88-196, [doi:10.1101/gr.6743907](https://doi.org/10.1101/gr.6743907).

9 Interval loci

The LocusPocus program ([AEGeAn Toolkit](#) version [53d092091c](#)) was used to compute interval loci (iLoci) from the genome annotation. The xtractore program was used to extract the iLocus sequences from the genome sequence. The iLocus sequences and annotations have been deposited in the Pdom Data Store at `/r1.2/interval-loci`.

9.1 Procedure (interactive)

We first need the genome sequence and corresponding annotations.

```
PdomData=/iplant/home/standage/Polistes_dominula/r1.2
iget ${PdomData}/genome-annotation/pdom-annot-r1.2.gff3
iget ${PdomData}/genome-assembly/pdom-scaffolds-unmasked-r1.2.fa.gz
gunzip pdom-scaffolds-unmasked-r1.2.fa.gz
```

Then we compute the iLocus coordinates, excluding gene-less iLoci if they are at the end of a scaffold.

```
locuspocus --intloci --skipends \
--delta=500 --verbose \
--outfile=pdom-loci-r1.2.gff3 \
pdom-annot-r1.2.gff3
```

The xtractore program gives us the iLocus sequences.

```
xtractore --type=locus --outfile=pdom-loci-r1.2.fa \
pdom-loci-r1.2.gff3 \
pdom-scaffolds-unmasked-r1.2.fa
```

For some analyses, we want gene coordinates expressed in reference to the iLocus to which they belong, as opposed to the entire scaffold. We transformed the annotations to iLocus-based coordinates and created two files, one including tRNA genes and one without.

```
perl glb2lcl.pl < pdom-loci-r1.2.gff3 \
| gt gff3 -retainids -sort -tidy -o pdom-annot-r1.2-iloci.gff3

perl glb2lcl.pl < pdom-loci-r1.2.gff3 \
| grep -v -e trnascan -e PdomTRNA \
| gt gff3 -retainids -sort -tidy -o pdom-annot-r1.2-iloci-sanstrna.gff3
```

9.2 Procedure (automated)

The same procedure can also be run in batch mode using the following commands (in the `interval-loci` directory).

```
make
make clean
```

10 Novel genes

Two previous analyses produced potential sources of novel genes warranting further investigation.

- The first potential source of novel genes are gene models predicted by Maker that have no matches against animal proteins in the NCBI nr database. These are described in the genome annotation documentation and are referred to hereafter as *unmatched gene models*.
- The second potential source of novel genes are assembled transcripts that have no matches in protein databases but are conserved among the 3 *Polistes* species. These are described in the transcript assembly documentation and are referred to hereafter as *unmatched TSAs*.

Here, we look for interval loci (iLoci) that have support from both of these sources of evidence. Relevant data has been uploaded to the Pdom Data Store at `r1.2/novel-genes`.

10.1 Procedure (interactive)

First we need to download the relevant data.


```
PdomData=/iplant/home/standage/Polistes_dominula/r1.2
iget ${PdomData}/transcript-assembly/pdom-tsa-r1.2-unmatched-pep.txt
iget ${PdomData}/transcript-alignment/pdom-tsa-masked-filtered.gff3
iget ${PdomData}/interval-loci/pdom-loci-r1.2.gff3
iget ${PdomData}/interval-loci/pdom-loci-r1.2-mrnamap.txt
iget ${PdomData}/genome-annotation/pdom-r1.2-no-anm-hits.ids
iget ${PdomData}/transcript-assembly/pdom-tsa-r1.2.fa.gz
gunzip pdom-tsa-r1.2.fa.gz
```

Next, identify iLoci to which unmatched TSAs align.

```
# Pull out alignment coordinates for unmatched TSAs
python align-keep.py pdom-tsa-r1.2-unmatched-pep.txt \
    pdom-tsa-masked-filtered.gff3 \
    > pdom-tsa-aligned-unmatched.gff3

# Find the corresponding iLoci
bedtools intersect -a <(grep $'\tlocus\t' pdom-loci-r1.2.gff3) \
    -b pdom-tsa-aligned-unmatched.gff3 -wa -u \
    | perl -ne 'm/(PdomILCr1.2-\d+)/ and print "$1\n" ' \
    | sort | uniq > pdom-tsa-aligned-unmatched-loci.txt
```

Then look at unmatched gene models and determine the iLoci to which they correspond.

```
./selex --out=1 pdom-r1.2-no-anm-hits.ids pdom-loci-r1.2-mrnamap.txt \
    | sort | uniq > pdom-r1.2-no-anm-hits-iloci.txt
```

Finally, determine which iLoci appear in both lists (iLoci to which unmatched TSAs align and iLoci containing unmatched gene models).

```
comm -12 pdom-r1.2-no-anm-hits-iloci.txt pdom-tsa-aligned-unmatched-loci.txt \
    > pdom-loci-r1.2-pot-trgs.txt
```

10.2 Procedure (automated)

The same procedure can also be run in batch mode using the following commands (in the `novel-genes` directory).

```
make
make clean
```